


Everything you always wanted to know about storage in Kubernetes

(But were too afraid to ask)

Johannes Wagner, Hendrik Land

2022-11-22

© 2022 NetApp, Inc. All rights reserved.



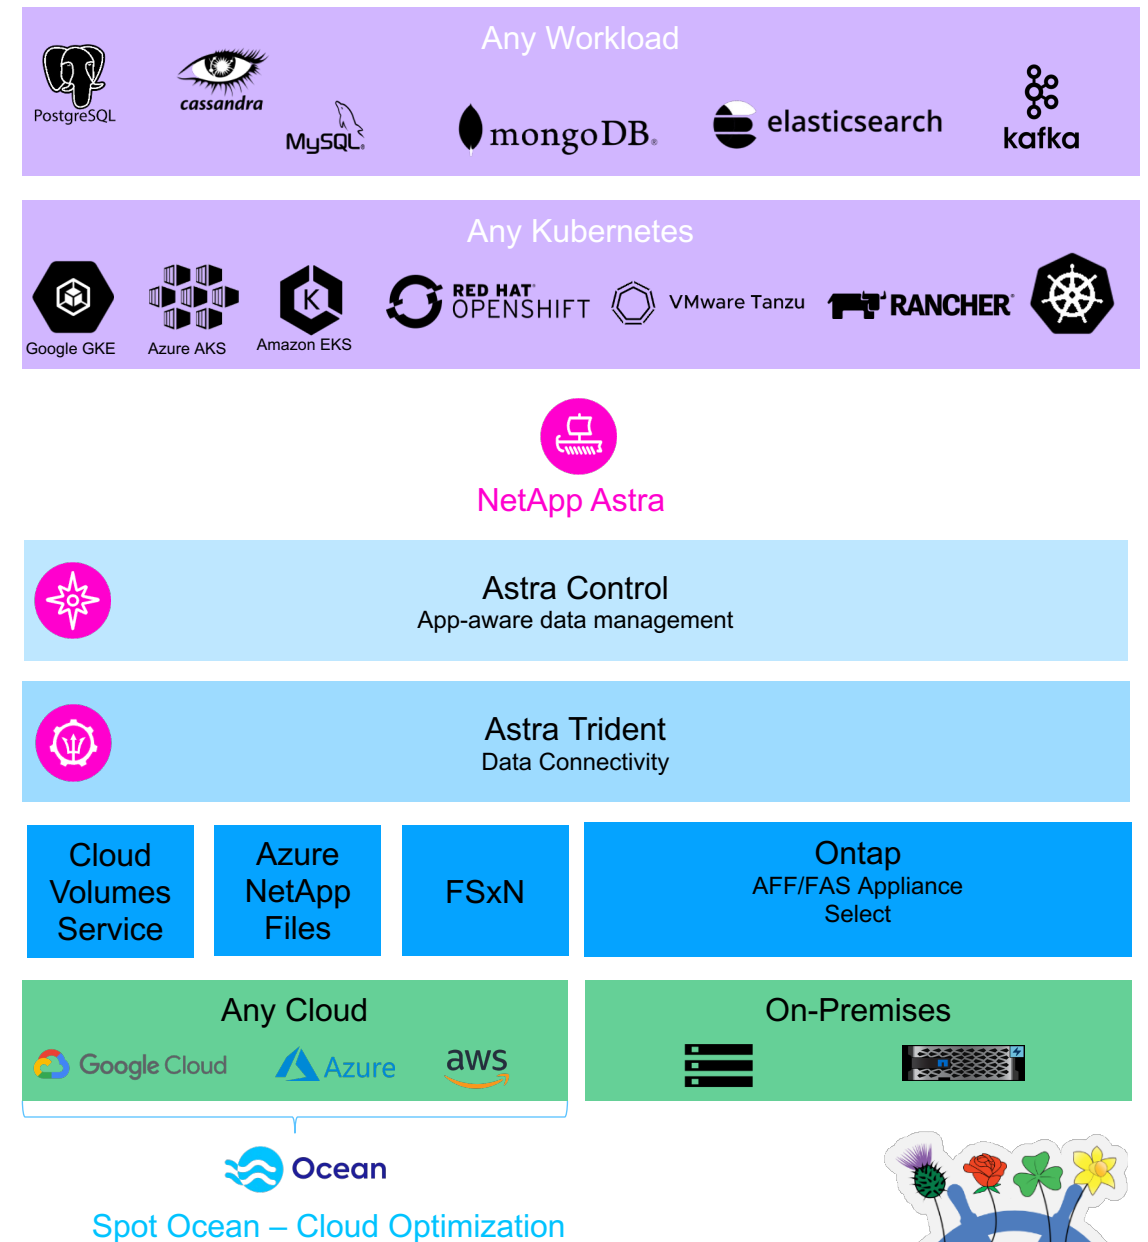
NetApp – Who are we?

- **Astra – Storage & Data Management for K8s**

- On-Premises & Any-Cloud
- Cloud-Native 1st party storage services
- Any K8s

- **Spot – Continuous Cloud Optimization**

- Ensure availability and optimize cost
- Container-driven autoscaling
- Container right-sizing



Storage in Kubernetes - Volume Types

Volume: Directory accessible to the containers in a pod

Resources exposed as a volume

- ConfigMap
- Secret
- ...

Storage – Place to store data


- In-tree & flexVolume storage drivers – deprecated!
- emptyDir
- Local
- PersistentVolumeClaim & CSI Drivers

CSI – Container Storage Interface

- Open Standard for attaching storage to containers based on file and block storage
- Provides dynamic provisioning of storage resources
- Provides storage as a file system to the container
- CSI alpha in K8s 1.9, beta in 1.11, GA in 1.13
- CNCF standard (though mostly used with K8s)

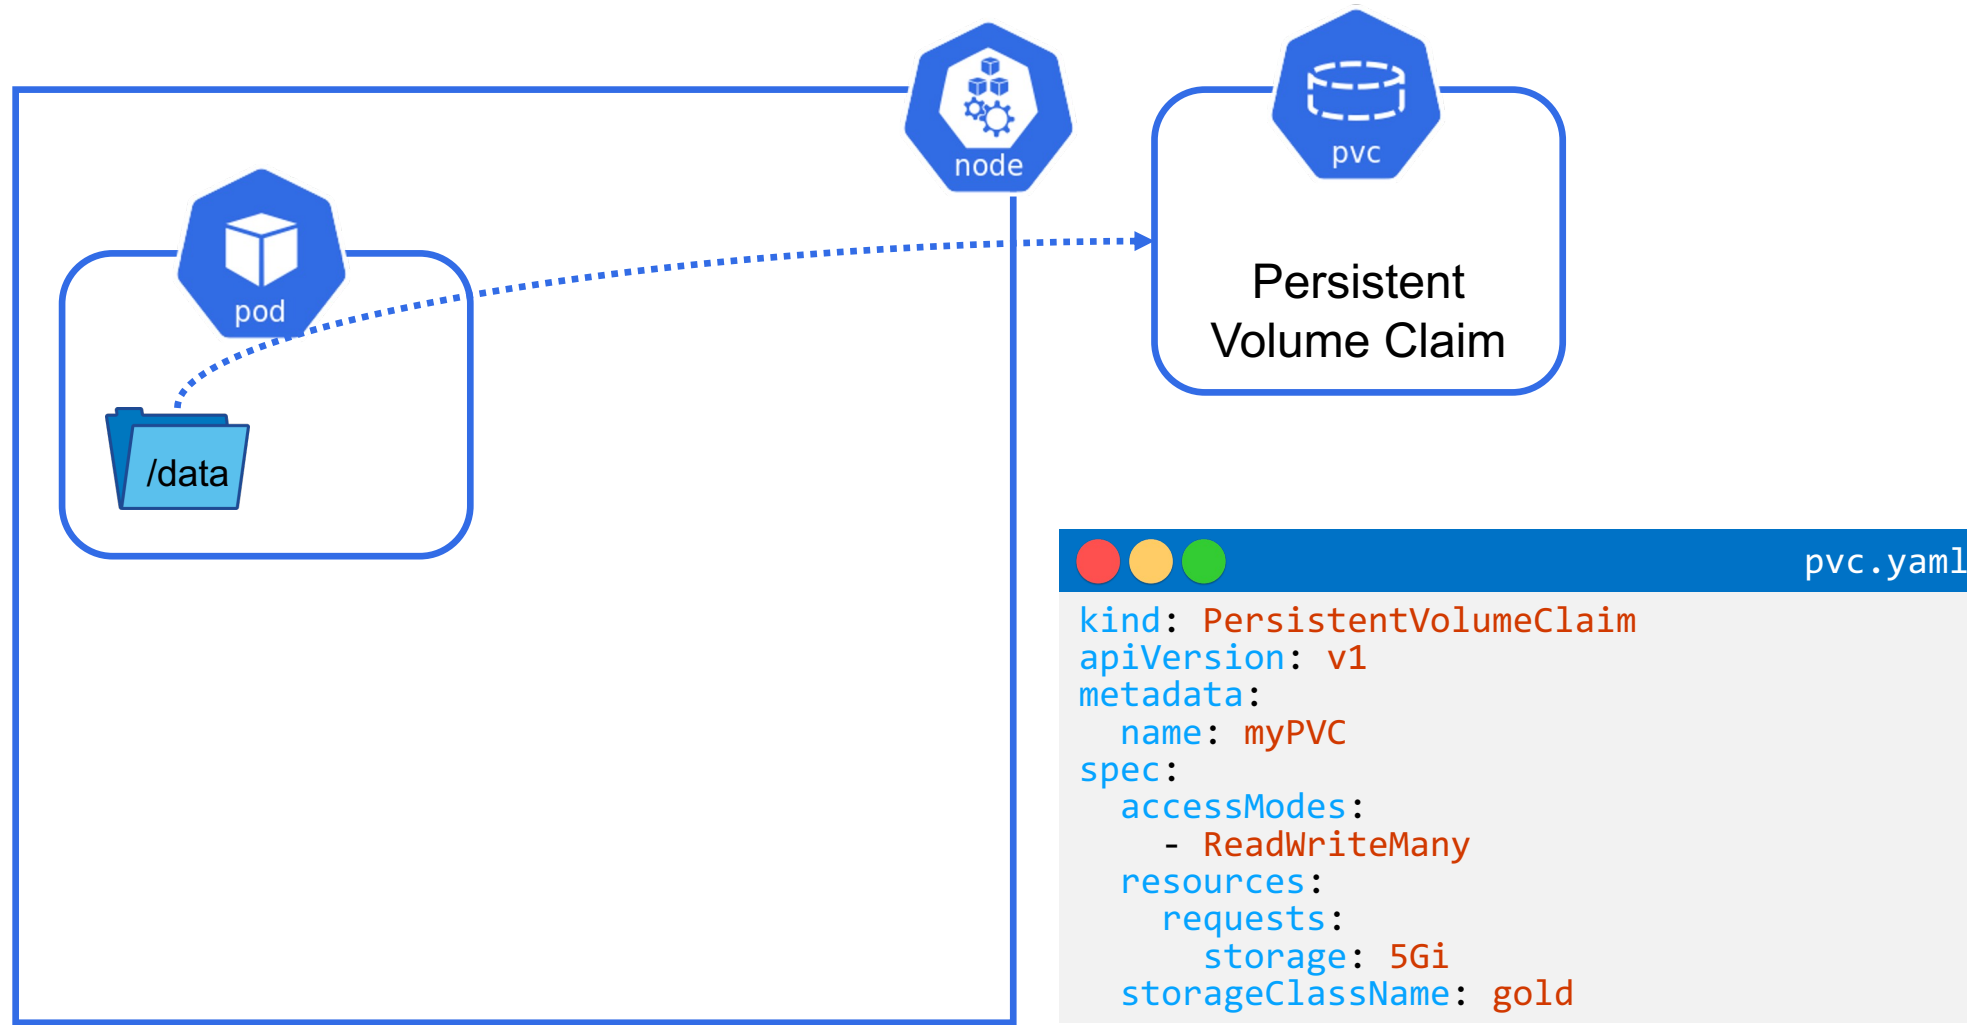


Dynamic storage provisioning in K8s

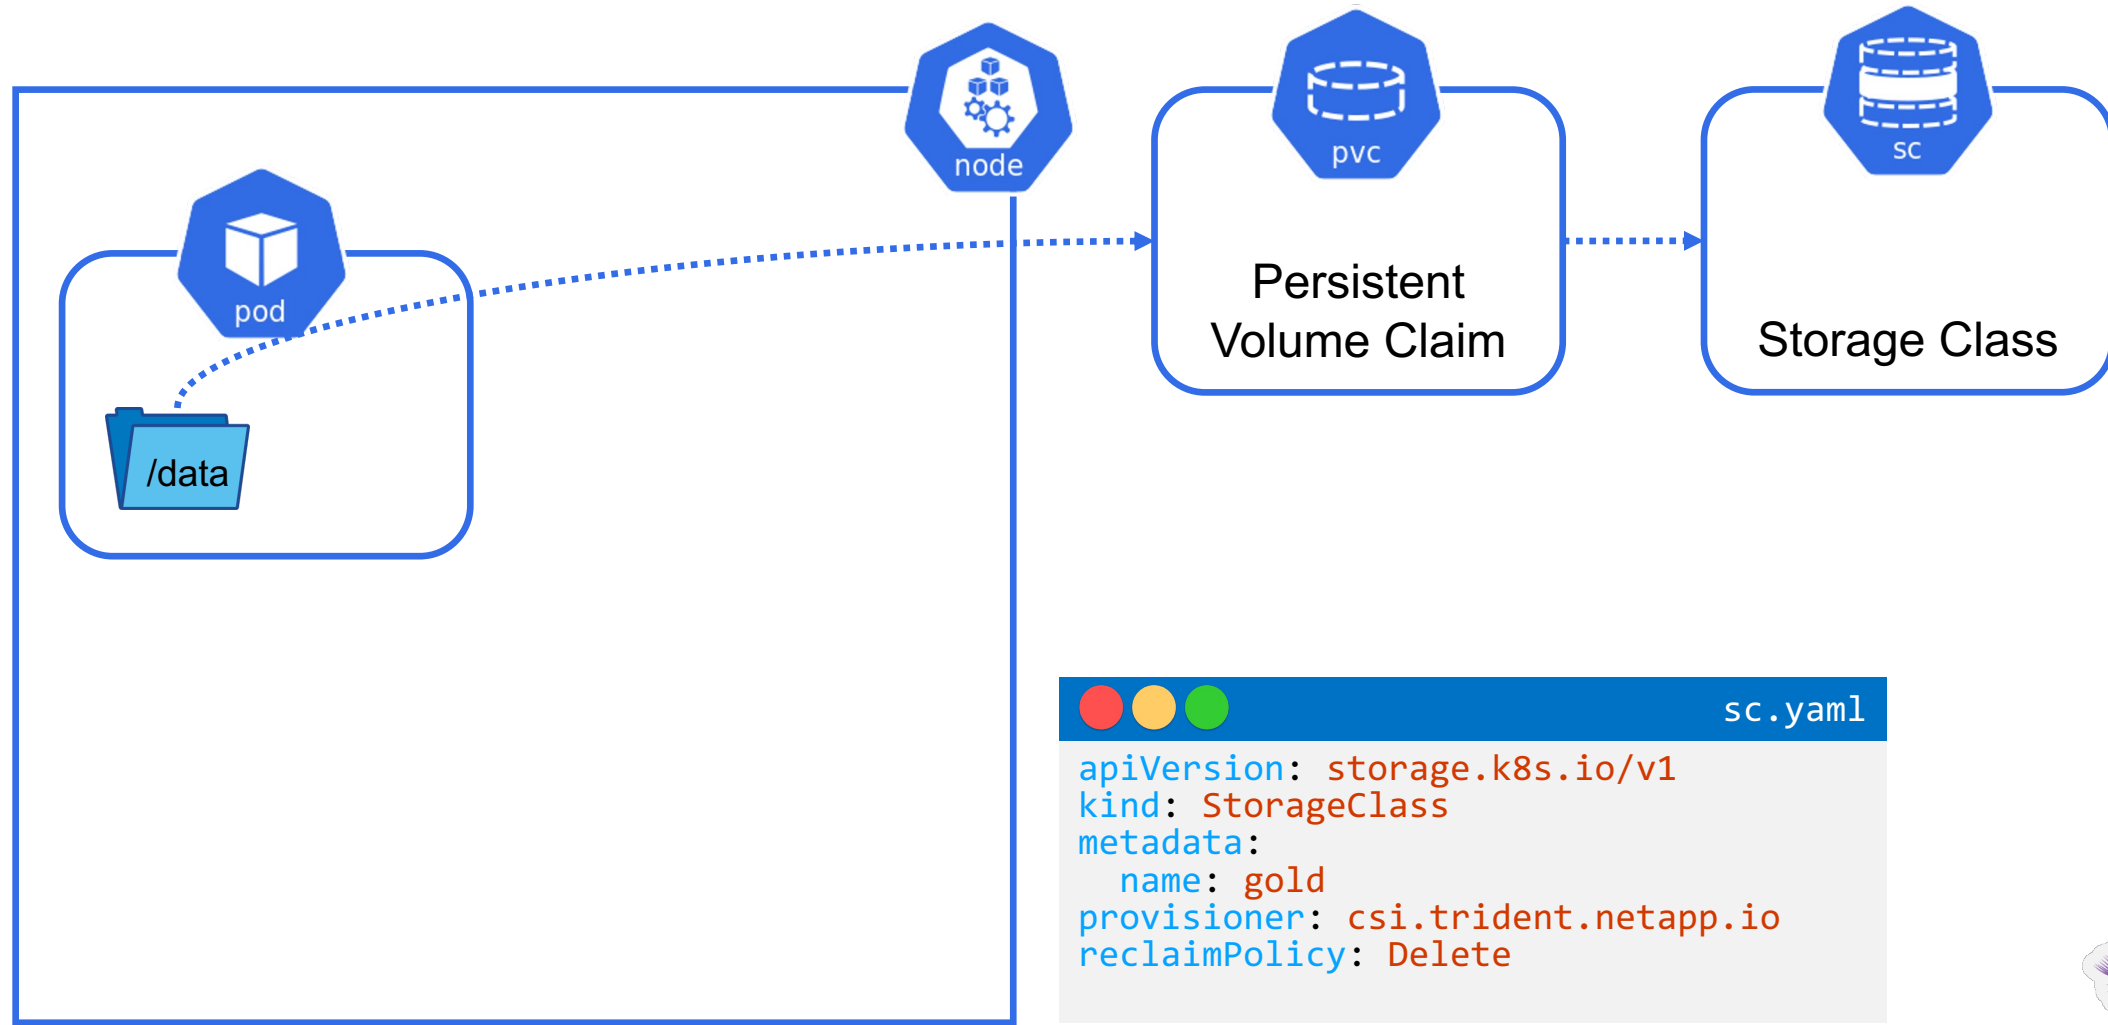


```
pod.yaml
kind: Pod
apiVersion: v1
metadata:
  name: myLittlePod
spec:
  volumes:
  - name: myVol
    persistentVolumeClaim:
      claimName: myPVC
  containers:
  - name: ubuntu-test
    image: ubuntu
    command: ["/<do something>"]
    volumeMounts:
    - mountPath: "/data"
      name: myVol
```

Dynamic storage provisioning in K8s

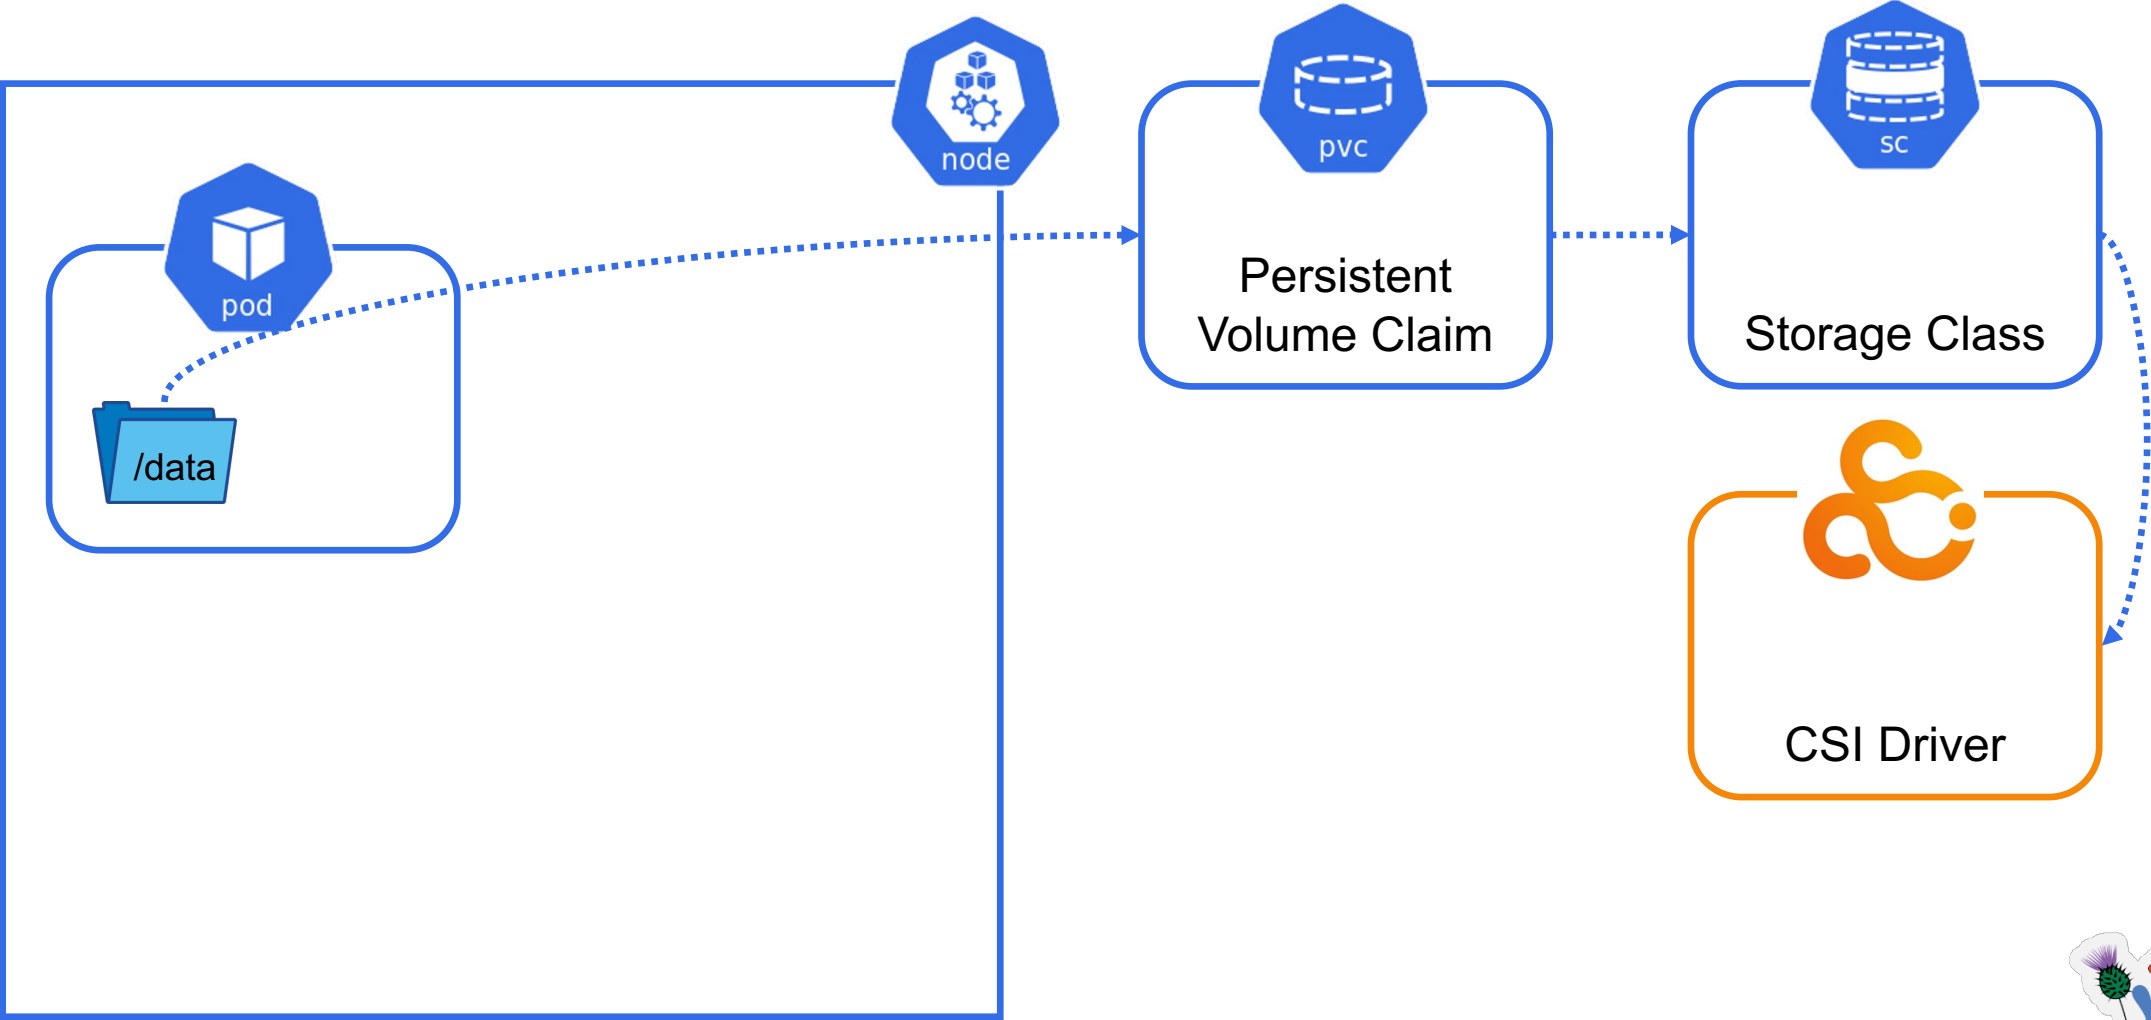


Dynamic storage provisioning in K8s

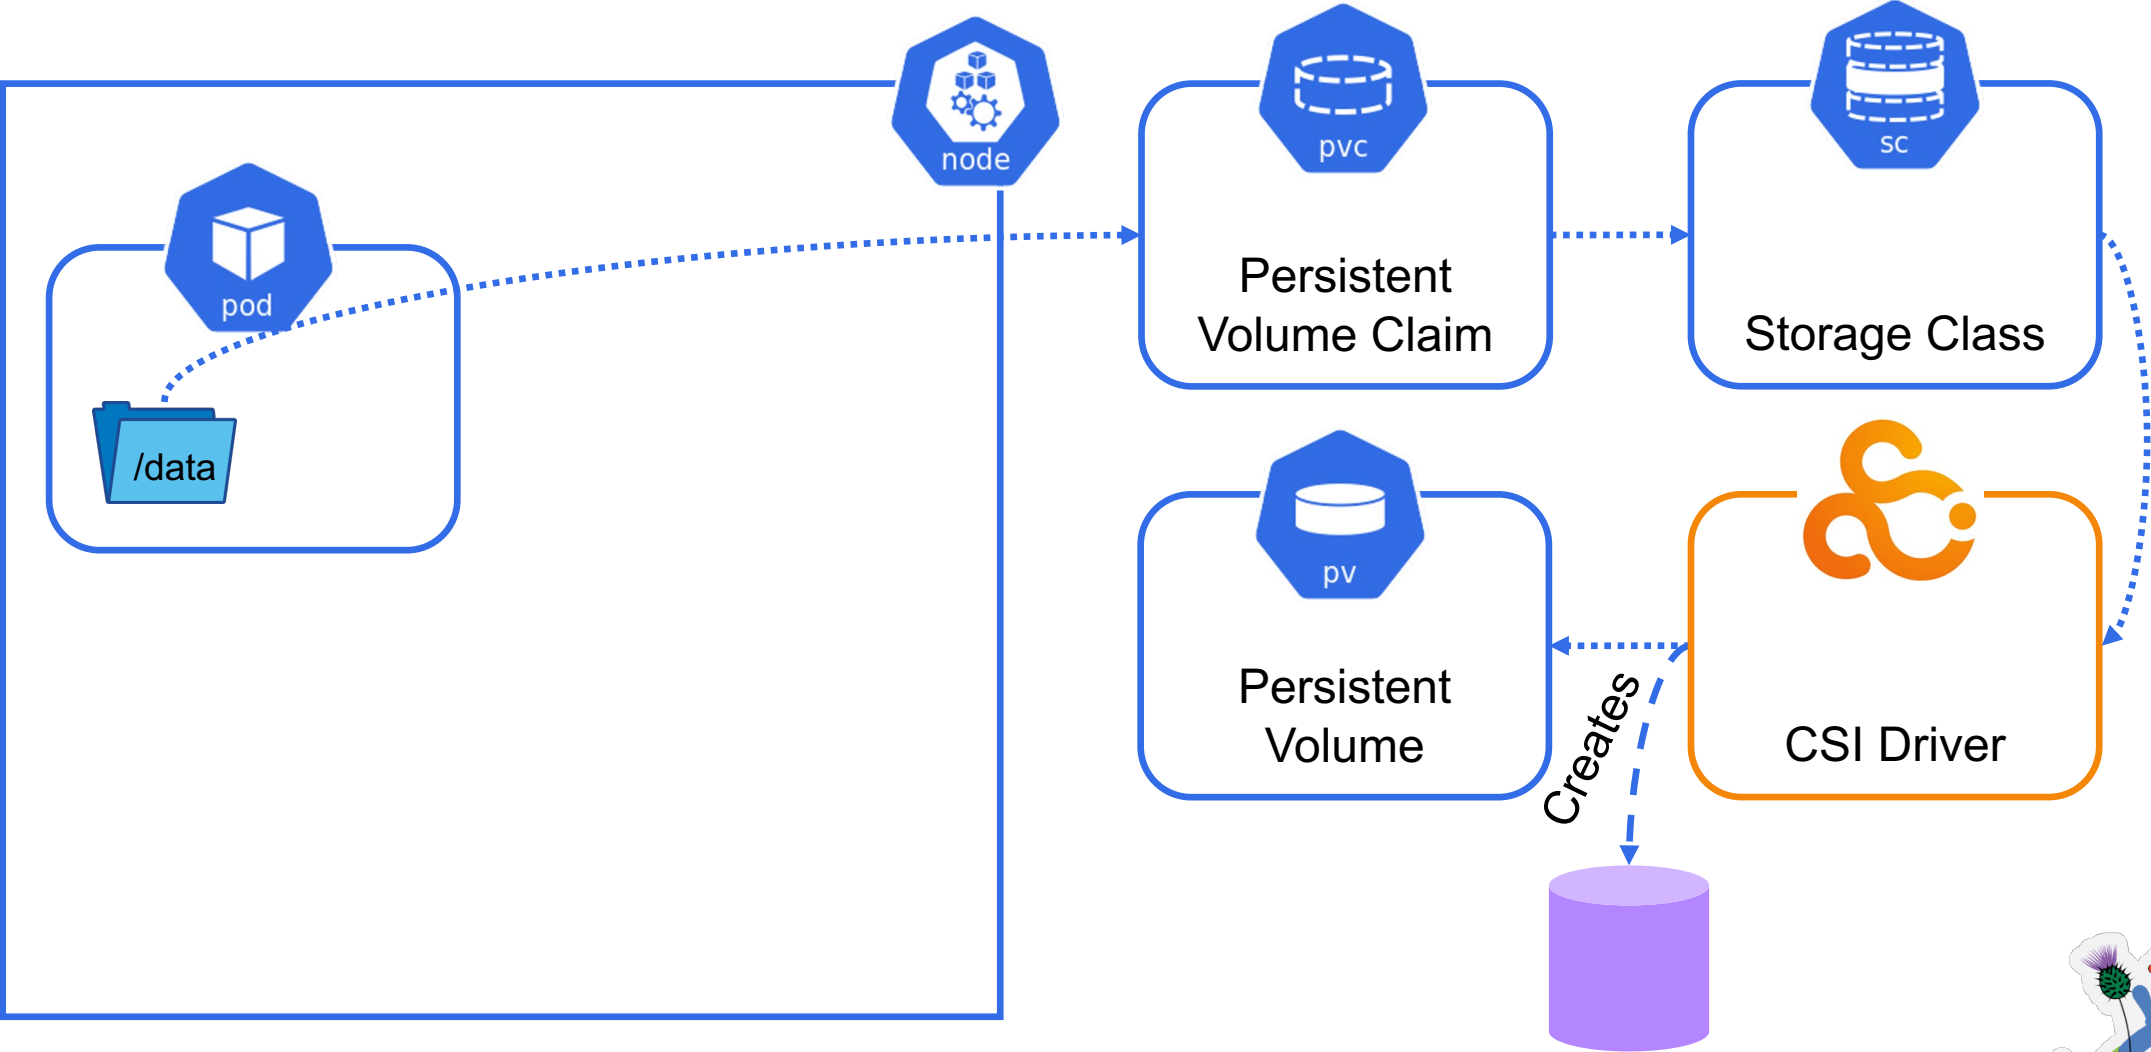


```
sc.yaml
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: gold
provisioner: csi.trident.netapp.io
reclaimPolicy: Delete
```

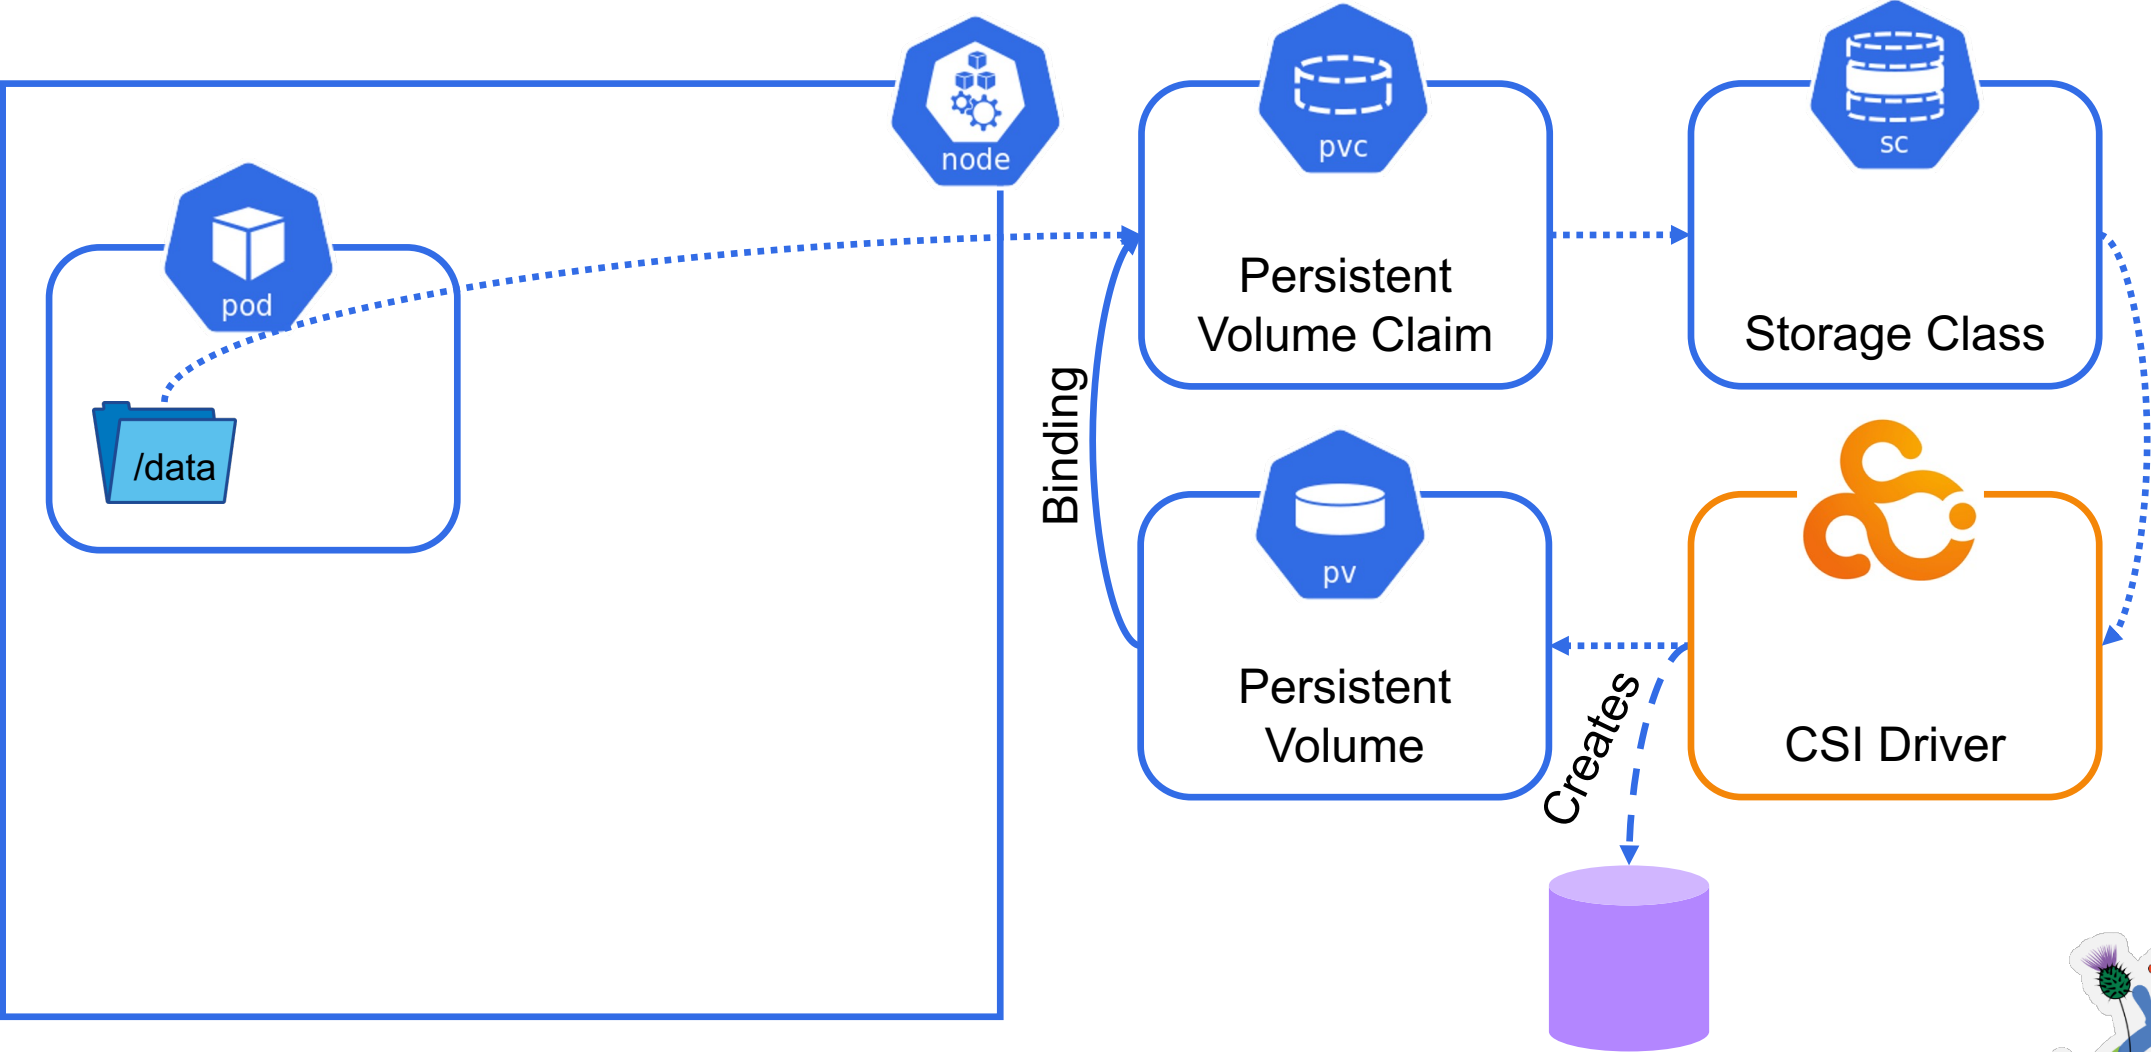
Dynamic storage provisioning in K8s



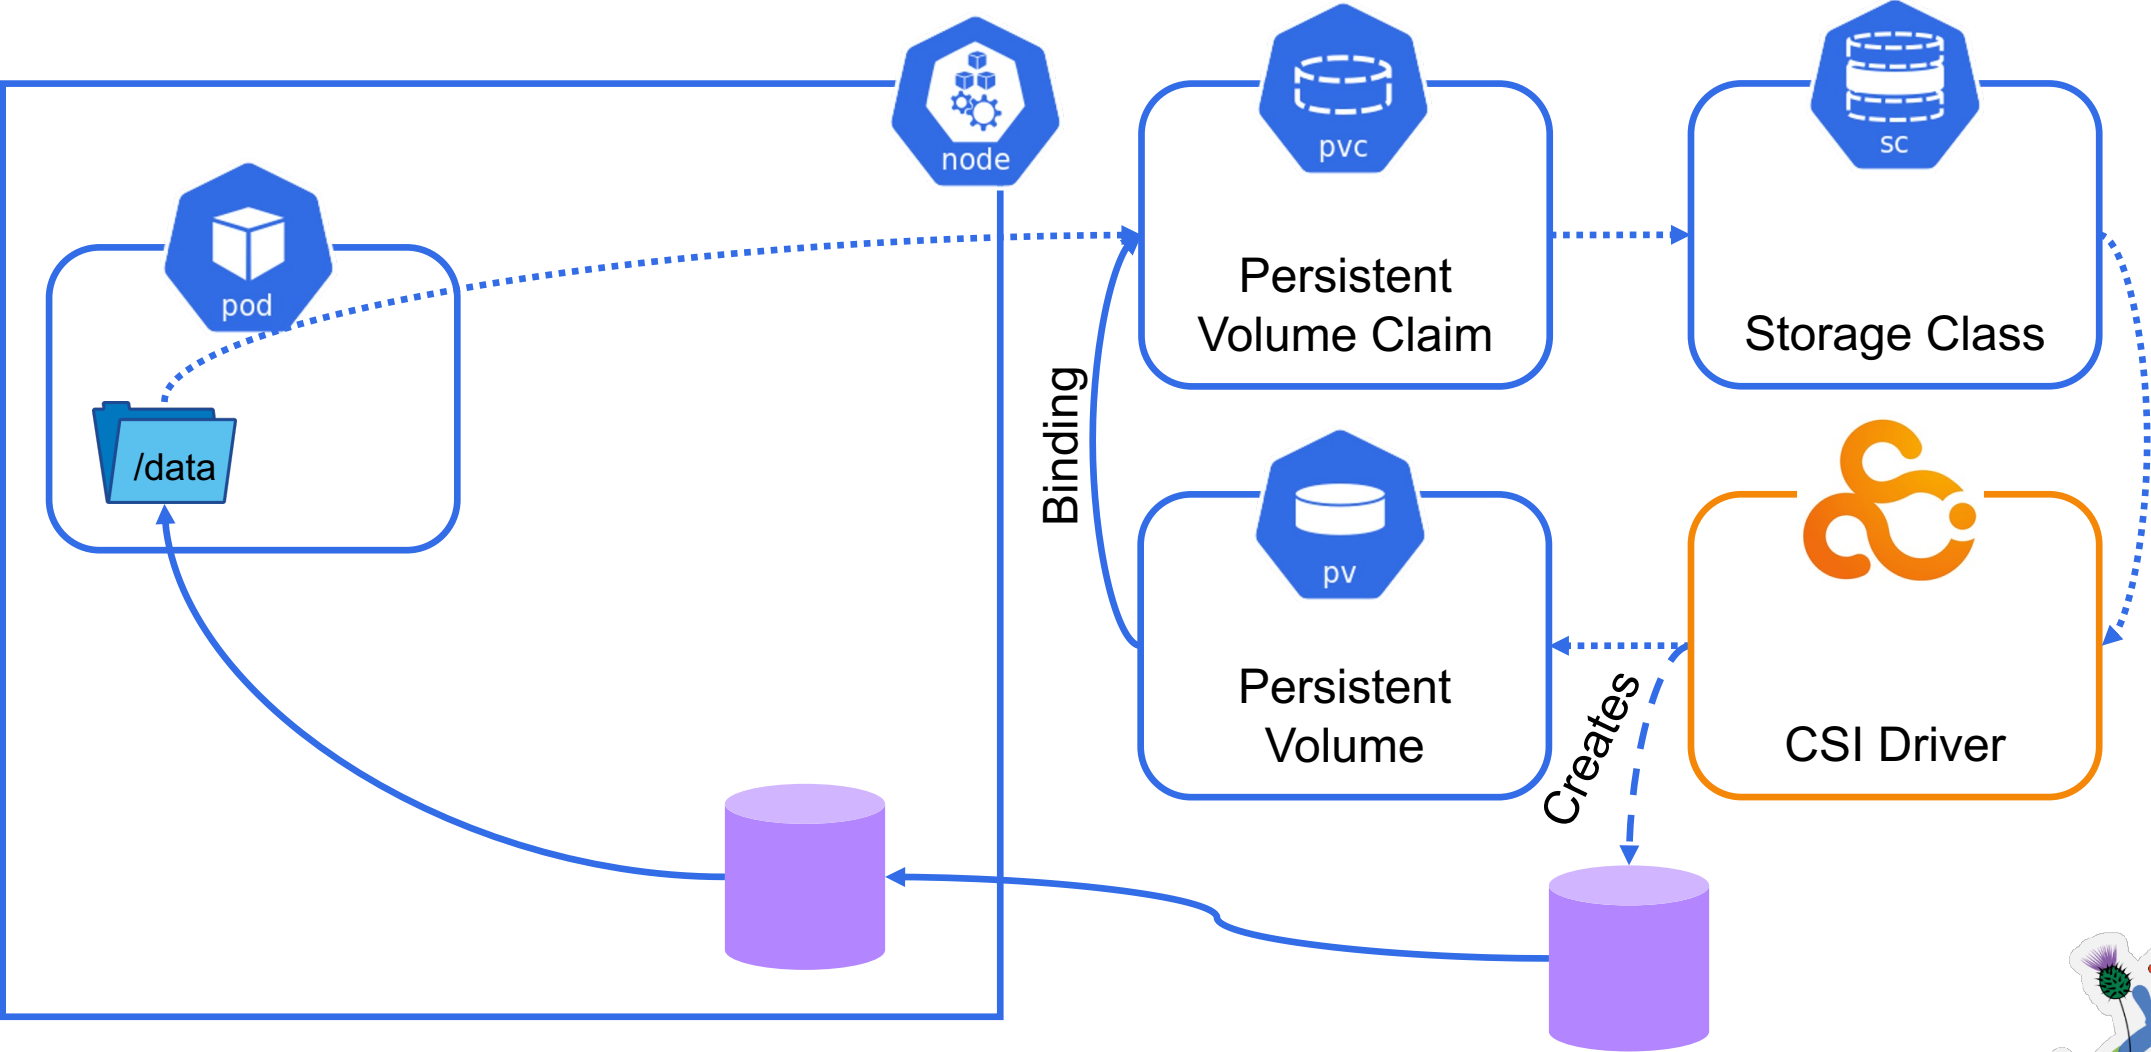
Dynamic storage provisioning in K8s



Dynamic storage provisioning in K8s



Dynamic storage provisioning in K8s



Persistent Volume Claim Access Modes

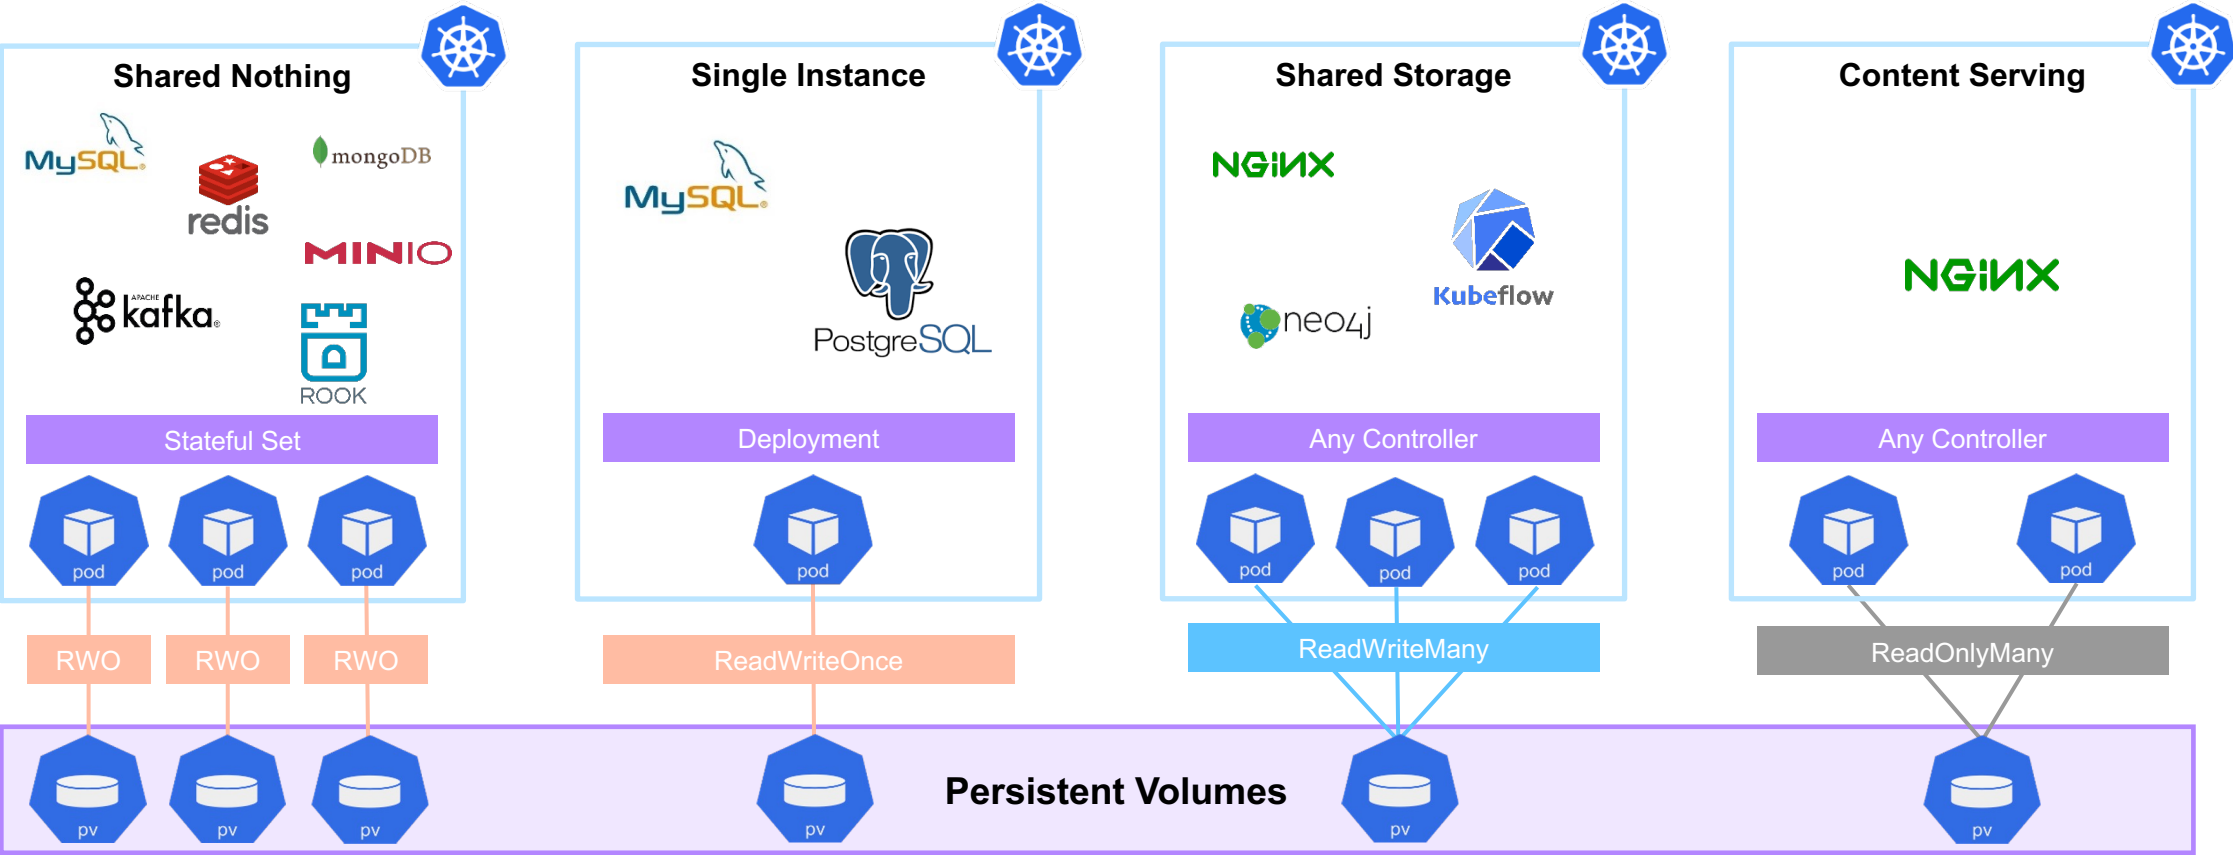
Different applications have different needs

- **ReadWriteOnce (RWO)**
 - Can be mounted read-write by a single node
- **ReadWriteMany (RWX)**
 - Can be mounted read-write by several nodes
- **ReadOnlyMany**
 - Read-Only access by several nodes
- **ReadWriteOncePod**
 - Read-Write access for a single pod only
 - New in K8s 1.22



Persistent Volume Claim Access Modes

Different applications have different needs



Storage Protocol Choices

File

- Shared filesystem (RWX)
- Good fit for Pod lifecycle
- Open Standard: NFS



Block

- Required by some applications (Prometheus, Kafka,...)
- Open Standards: iSCSI, NVMe



Object

- Data Service, not a file system inside the container
- “Standard”: S3




Container Storage Interface
(CSI)

Container Object Storage Interface
(COSI) – Alpha in K8s 1.25



Hands-on Part 1



Special
thanks to
Yves
Weisser

Before we start

- Please raise your hand or ask us if you have questions or problems
- Choose your username from the spreadsheet, enter your name to avoid double usage
- Pre-setup is important
- Scenario 01 contains basic setup steps and is a must, following scenarios depend on it
- Before Scenario 03 we want to do a short theory session again

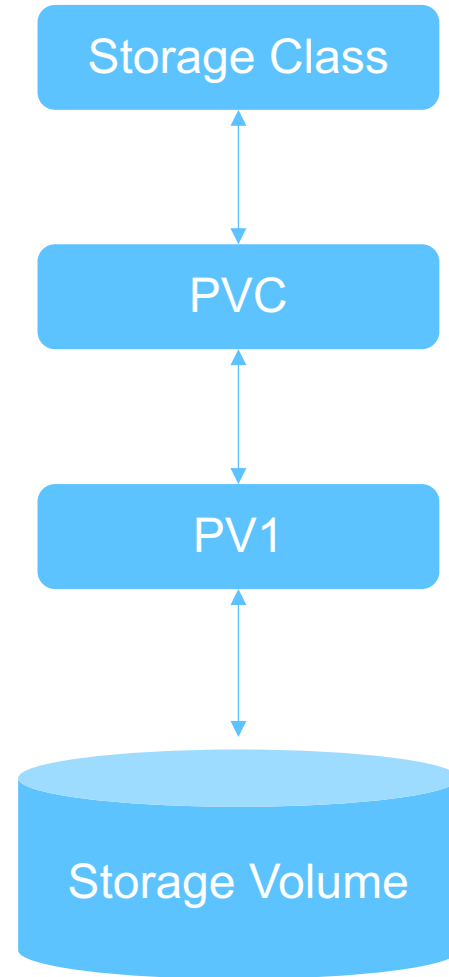


Hands-on Part 1

<https://github.com/kcdstorageews/lab>

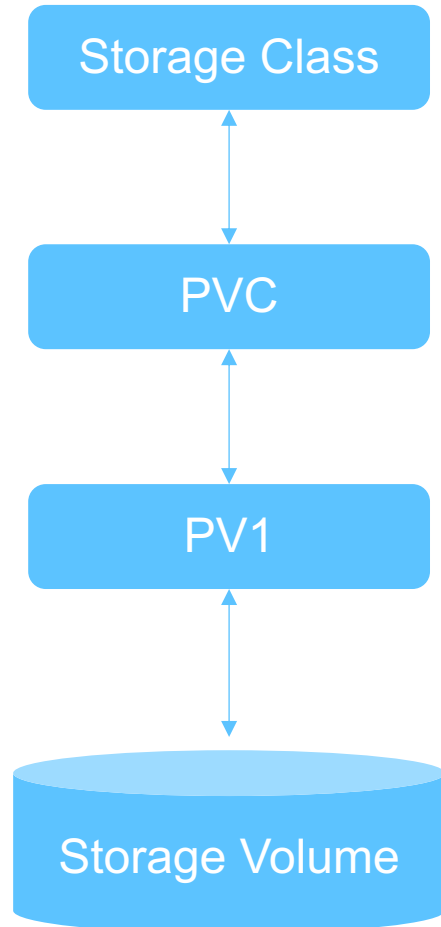
CSI Snapshots workflow

Initial configuration



CSI Snapshots workflow

Volume Snapshot Class creation



Volume
Snapshot Class



volumesnapshotclass.yaml

```
apiVersion: snapshot.storage.k8s.io/v1beta1
kind: VolumeSnapshotClass
metadata:
  name: csi-snapclass
driver: csi.trident.netapp.io
deletionPolicy: Delete
```

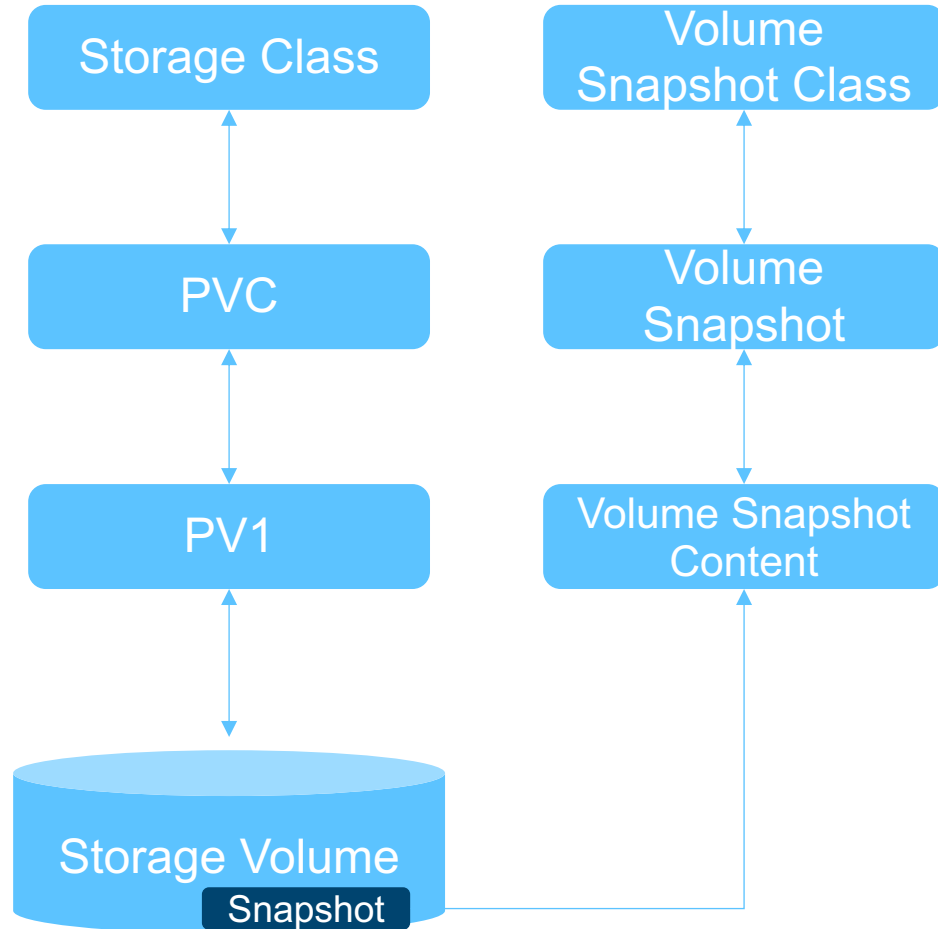
NOTE

the *deletionPolicy* parameter can be set to:

- Delete
- Retain

CSI Snapshots workflow

Volume Snapshot creation

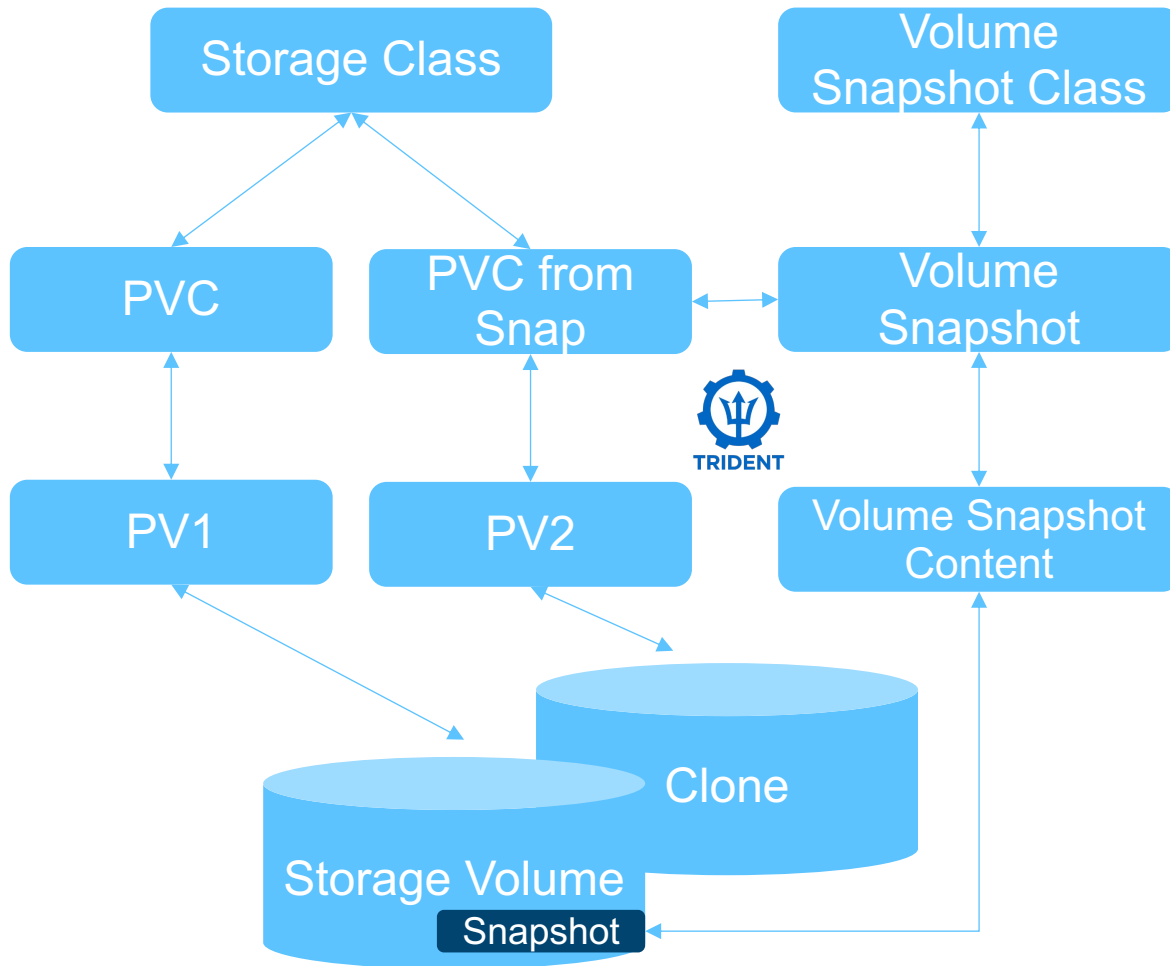


volume-snapshot.yaml

```
apiVersion: snapshot.storage.k8s.io/v1beta1
kind: VolumeSnapshot
metadata:
  name: volume-snapshot
spec:
  volumeSnapshotClassName: csi-snapclass
  source:
    persistentVolumeClaimName: pvc
```

CSI Snapshots workflow

Clone creation



```
pvc-from-snapshot.yaml

kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: pvc-from-snap
spec:
  accessModes:
    - ReadWriteMany
  resources:
    requests:
      storage: 1Gi
  storageClassName: sc-nas
  dataSource:
    name: volume-snapshot
    kind: VolumeSnapshot
    apiGroup: snapshot.storage.k8s.io
```

Consumption Control

- Standard K8s mechanisms such as [ResourceQuota](#) and [LimitRange](#) apply to storage as well
- Restrict capacity and number of PVCs per namespace with a [ResourceQuota](#)
 - Total capacity
 - Capacity per [StorageClass](#)
 - Total number of PVCs
 - Number of PVCs per [StorageClass](#)
- Remember: [StorageClass](#) is a global resource in the cluster
 - But you can assign a 0 byte quota
- Define Minimum and Maximum size of an individual PVC with [LimitRange](#)

Generic Ephemeral Volumes

- K8s provides different types of ephemeral volumes
 - emptyDir
 - CSI ephemeral volumes
 - **generic ephemeral volumes**
- Generic ephemeral volumes provide additional features (beta in K8s 1.21, GA in 1.23)
 - External storage to not exceed disk capacity of local node
 - Snapshot & Cloning Capability
 - Resizing
 - StorageClass
- Set StorageClass binding mode to [WaitForFirstConsumer](#) so K8s scheduler is free to choose any node




Hands-on Part 2

<https://github.com/kcdstorageews/lab>

EVERYTHING you always wanted to know about storage in Kubernetes?

OK, there's more...



- CSI Topology
- (Capacity) monitoring
- Non-graceful shutdown
- Security

